A_Memorix: An API-First Hybrid Memory Service for Temporal and Graph-Aware Retrieval (ACL 2026 Demo)

Chen Xi
Independent Researcher
China
`https://github.com/A-Dawn`

Abstract

We present **A_Memorix**, an API-first memory service that unifies vector retrieval, graph retrieval, temporal filtering, and online memory maintenance in a standalone runtime. The system targets long-horizon conversational and agentic scenarios where knowledge must be continuously ingested, retrieved under semantic and temporal constraints, and revised over time. A_Memorix combines dual-path retrieval (paragraph and relation paths), sparse fallback with rank fusion, minute-level temporal querying, and lifecycle-oriented memory operations including protect, reinforce, freeze, prune, and restore. The demo exposes both a new `/v1/*` API and a backward-compatible `/api/*` layer, with a web visualization entry point. In local measurements, the metadata temporal path reaches 826.8 QPS at 5k paragraphs (p95 1.515 ms), while HTTP end-to-end `/v1/query/time` reaches 121.2 QPS at 5k paragraphs (p95 26.166 ms). Artifact links required for the demo track are provided in Section 6.

1 Introduction

LLM applications with persistent memory require more than a standalone retriever: they also require memory governance, temporal constraints, and operationally safe update and recovery semantics. In practice, these capabilities are frequently distributed across separate tools, which increases integration complexity. Retrieval-augmented generation systems further highlight the central role of robust external memory interfaces in knowledge-intensive settings (Lewis et al., 2020).

In this demo paper, we present A_Memorix as a single-process, API-centric memory service with:

- dual-path retrieval over vectors and graph relations;

- temporal retrieval with structured minute-level query formats;

- memory lifecycle operations (protect/reinforce/freeze/prune/restore);

- asynchronous import and transcript-summary task orchestration.

Intended users include applied NLP engineers, agent platform developers, and research teams that require self-hosted long-term memory APIs with temporal and governance controls. Source code and reproducibility artifacts are available at `https://github.com/A-Dawn/A_memorix/releases/tag/acl2026-demo-v1.1.0`.

2 System Overview

2.1 Runtime Architecture

A_Memorix is implemented as a FastAPI service with a shared runtime context and three persistent stores:

- **Vector store** (`data/vectors`): FAISS-based append-only storage with SQ8 quantization and fallback flat index (Johnson et al., 2021).

- **Graph store** (`data/graph`): sparse adjacency matrix with relation-hash mapping.

- **Metadata store** (`data/metadata`): SQLite metadata for paragraphs, entities, relations, FTS, async tasks, transcript sessions, and recycle-bin records.

2.2 API Surfaces

The service exposes two API layers:

- `/v1/*` for new integrations;

- `/api/*` for backward compatibility with historical contracts.

This dual interface supports migration without disrupting existing clients.

Case	Paras	QPS	p95	p99
s_seed7	1200	3046.8	0.454	0.560
m_seed42	5000	826.8	1.515	1.670
m_noperson	5000	1081.2	1.275	1.328
l_seed42	10000	425.6	2.976	3.225

Table 1: Metadata-path temporal micro-benchmark (latencies in ms).

Case	Paras	QPS	p95	p99
e2e_s_seed7	800	156.1	24.160	27.345
e2e_m_seed42	3000	132.9	25.086	27.529
e2e_m_noperson	3000	138.9	23.149	28.898
e2e_l_seed42	5000	121.2	26.166	29.620

Table 2: HTTP end-to-end benchmark for `/v1/query/time` (latencies in ms).

2.3 Positioning Versus Existing Tooling

A_Memorix is positioned between generic vector databases and tightly coupled agent frameworks. Relative to vector-only memory layers, it provides graph relations, temporal filters, and explicit lifecycle controls within a single service boundary. Relative to framework-specific memory modules, it offers framework-agnostic HTTP interfaces (`/v1/*`, `/api/*`) and reproducible standalone deployment paths that can be reused across heterogeneous application stacks.

2.4 Core Retrieval and Memory Pipeline

Dual-path retrieval combines paragraph search and relation-centric search. Sparse BM25 fallback (FTS5-based) follows probabilistic retrieval principles (Robertson and Walker, 1994) and can be merged with dense candidates via weighted RRF fusion (Cormack et al., 2009). Optional Personalized PageRank reranking adds entity-aware graph priors (Page et al., 1999).

Temporal filtering uses interval-overlap semantics over `event_time`, `event_time_start`, and `event_time_end`, with optional `created_at` fallback. Query-time formats are intentionally strict to reduce ambiguity:

- `YYYY/MM/DD`

- `YYYY/MM/DD HH:mm`

Memory maintenance periodically decays edge weights, freezes low-weight unprotected relations, prunes expired inactive relations into a recycle bin, and supports explicit restore operations.

3 Benchmark and Test Status

3.1 Metadata Temporal Micro-Benchmark

We benchmarked the metadata temporal query path (`query_paragraphs_temporal`) on synthetic data.

3.2 `/v1/query/time` End-to-End Benchmark

We additionally ran HTTP E2E measurements by synthesizing temporal data, launching local service instances, and issuing POST requests to `/v1/query/time`.

3.3 In-Tree Tests

The repository includes in-tree tests for time parsing, temporal filtering, summary import, and search execution behavior. Current status:

- `python -m pytest -q` → 12 passed.

3.4 Human Evaluation Status

We have not yet completed a formal human-subject study for this demo paper. Current evidence is system-oriented, including deterministic in-tree tests, micro-benchmarks, and HTTP end-to-end benchmarks. We plan to add task-level user evaluation in a later release.

4 Related Work

Retrieval-augmented generation systems demonstrate that external memory access is central to knowledge-intensive NLP workflows (Lewis et al., 2020). In this landscape, dense vector retrieval infrastructure (Johnson et al., 2021), sparse probabilistic retrieval (Robertson and Walker, 1994), rank-fusion methods (Cormack et al., 2009), and graph-centrality signals (Page et al., 1999) are often studied separately. A_Memorix integrates these components into one operational service and emphasizes temporal query semantics and memory lifecycle governance in addition to retrieval quality.

5 Demo Walkthrough

The live demo sequence is:

1. start service and verify `/healthz`, `/readyz`;

2. create import task (`/v1/import/tasks`) and poll status;

3. issue temporal search (`/v1/query/time`) with request fields such as `query`, `time_from`, `time_to`, and `top_k`;

4. apply memory actions (`/v1/memory/protect`, `/v1/memory/reinforce`, `/v1/memory/restore`);

5. run profile and summary APIs (`/v1/person/*`, `/v1/summary/tasks`).

6 Artifacts and Access

To satisfy ACL 2026 Demo artifact expectations, we provide:

- Source code and installable package: GitHub release (acl2026-demo-v1.1.0); package asset: acl2026-demo-v1.1.0.zip

- Reproducible Docker entry: repository `Dockerfile` and benchmark scripts

- License: `AGPL-3.0` (same as repository root license)

- Demo video ($\leq$2.5 min): video.mp4 (release asset)

7 Limitations and Ethics

Limitations. Our benchmark workloads are synthetic and should not be interpreted as full production SLOs. Retrieval quality metrics still depend on configured external OpenAI-compatible providers.

Ethics. The system stores user text and derived relational memory; deployments should enforce access control, retention policies, and user-visible deletion/recovery controls. Person profile aggregation should be opt-in and policy-governed.

8 Conclusion

A_Memorix provides a practical, API-first memory system that combines retrieval, temporal reasoning, and memory lifecycle management in a deployable standalone service. Our demo and benchmarks show operational viability for ACL system demonstration use cases and provide a reproducible baseline for future extensions.

References

Gordon V. Cormack, Charles L. A. Clarke, and Stefan Büttcher. 2009. Reciprocal rank fusion outperforms condorcet and individual rank learning methods. In *Proceedings of the 32nd Annual International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '09)*, pages 758–759. ACM.

Jeff Johnson, Matthijs Douze, and Hervé Jégou. 2021. Billion-scale similarity search with gpus. *IEEE Transactions on Big Data*, 7(3):535–547.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems 33 (NeurIPS 2020)*. Curran Associates, Inc.

Lawrence Page, Sergey Brin, Rajeev Motwani, and Terry Winograd. 1999. The PageRank citation ranking: Bringing order to the web. Technical Report SIDL-WP-1999-0120, Stanford InfoLab.

Stephen E. Robertson and Steve Walker. 1994. Some simple effective approximations to the 2-poisson model for probabilistic weighted retrieval. In *Proceedings of the 17th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '94)*, pages 232–241. ACM/Springer.
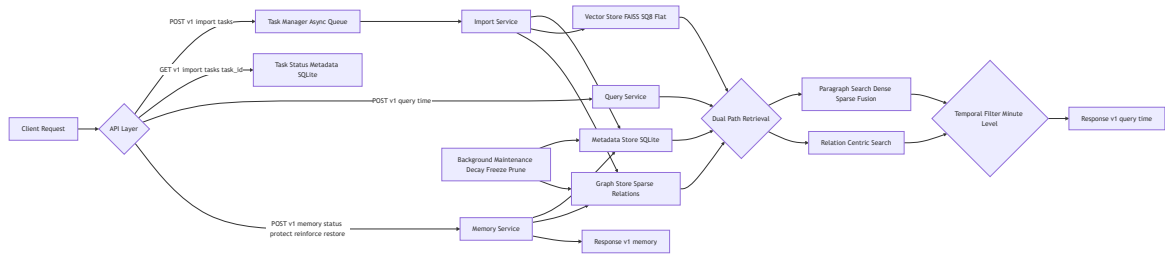
Figure 1: System flow used in the ACL demo. Import tasks, temporal query, and memory APIs are modeled as separate request paths over shared stores; freeze/prune are background maintenance operations rather than direct `/v1/memory/*` endpoints.